

Statistical Learning Lab

Assignment - 4

Cross-validation and Bootstrapping

**Submitted by,
Ben Abraham Biju
22IM10048**

1. Load the dataset “manufacturing.csv”. Display the first few rows of the dataset. Take “Quality Rating” as the response variable.

The dataset was loaded using the read.csv() function and stored in the variable data.

```
> data <- read.csv("C:/Users/benab/OneDrive - iitkgp.ac.in/Desktop/Sem 6/SL Lab/Lab 2/manufacturing.csv", header = TRUE, stringsAsFactors = FALSE)
> head(data, 10)
  Temperature...C. Pressure..kPa. Temperature.x.Pressure Material.Fusion.Metric Material.Transformation.Metric Quality.Rating
1      209.7627      8.050855      1688.769      44522.22      9229576      99.99997
2      243.0379     15.812068      3842.931      63020.76     14355367     99.98570
3      220.5527      7.843130     1729.823     49125.95     10728389     99.99976
4      208.9766     23.786089     4970.737     57128.88     9125702     99.99997
5      184.7310     15.797812     2918.345     38068.20     6303792    100.00000
6      229.1788      8.498306     1947.632     53136.69     12037072     99.99879
7      187.5174     19.412851     3640.248     42478.69     6593260    100.00000
8      278.3546      7.070944     1968.230     77834.82     21567222     95.73272
9      292.7326     20.432896     5981.374     94223.15     25084522     64.62360
10     176.6883     14.145782     2499.394     34049.37     5515789    100.00000
> |
```

2. Fit polynomial models between Quality ~ Temp. Vary the degree of polynomial on temperature from 1 to 5 (temp, temp^2, temp^3 etc.). Perform LOOCV, k-fold CV for k=5 and 10 and compare the cross-validation MSE errors for different degrees of polynomials. Create a table showing the CV errors for different degrees of polynomials and for different CV techniques. Plot the results. Discuss which degree of polynomial is preferable.

The range of degrees of the polynomial is set from 1 to 5, and a null list is made for each type of cross validation method. The dataset is stored in the variable data sample.

```
data_sample <- data

degrees <- 1:5
cv.error.loocv <- rep(0, length(degrees))
cv.error.5fold <- rep(0, length(degrees))
cv.error.10fold <- rep(0, length(degrees))

for (i in degrees) {
  start_time <- Sys.time()
  cat("Starting polynomial degree:", i, "\n")

  glm.fit <- glm(Quality ~ poly(Temperature, i, raw = TRUE), data = data_sample)

  # LOOCV
  cv.loocv <- cv.glm(data_sample, glm.fit, K = nrow(data_sample))
  cv.error.loocv[i] <- cv.loocv$delta[1]

  # 5-fold CV
  cv.5 <- cv.glm(data_sample, glm.fit, K = 5)
  cv.error.5fold[i] <- cv.5$delta[1]

  # 10-fold CV
  cv.10 <- cv.glm(data_sample, glm.fit, K = 10)
  cv.error.10fold[i] <- cv.10$delta[1]

  end_time <- Sys.time()
  cat("Finished degree:", i, "in", round(difftime(end_time, start_time, units = "secs"), 2), "seconds\n")
}

Starting polynomial degree: 1
Finished degree: 1 in 20.3 seconds
Starting polynomial degree: 2
Finished degree: 2 in 21.97 seconds
Starting polynomial degree: 3
Finished degree: 3 in 25.13 seconds
Starting polynomial degree: 4
Finished degree: 4 in 28.43 seconds
Starting polynomial degree: 5
Finished degree: 5 in 32.67 seconds
```

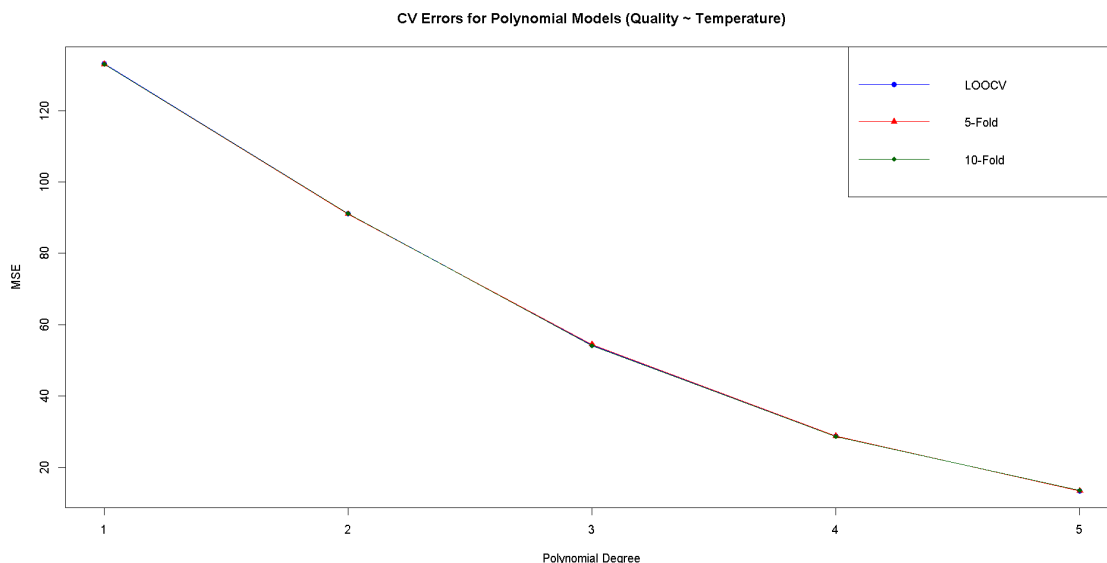
Once the loop has been executed for all degrees, the errors are stored as a dataframe and printed.

```
poly.cv.results <- data.frame(
  Degree = degrees,
  LOOCV_MSE = cv.error.loocv,
  CV5_MSE = cv.error.5fold,
  CV10_MSE = cv.error.10fold
)

[1] "CV Errors for Polynomial Models (Quality ~ Temperature):"
> print(poly.cv.results)
  Degree LOOCV_MSE  CV5_MSE  CV10_MSE
1      1 133.07880 132.88638 133.36005
2      2  91.19322  91.26217  91.21856
3      3  54.23607  54.20037  54.57484
4      4  28.78949  28.64342  28.79402
5      5  13.52725  13.46535  13.47917
```

The errors for different degrees of polynomial and different methods are plotted

```
plot(degrees, cv.error.loocv, type = "o", col = "blue", pch = 16,
     xlab = "Polynomial Degree", ylab = "MSE",
     main = "CV Errors for Polynomial Models (Quality ~ Temperature)")
lines(degrees, cv.error.5fold, type = "o", col = "red", pch = 17)
lines(degrees, cv.error.10fold, type = "o", col = "darkgreen", pch = 18)
legend("topright", legend = c("LOOCV", "5-Fold", "10-Fold"),
     col = c("blue", "red", "darkgreen"), pch = c(16, 17, 18), lty = 1)
```



The results show that the least MSE occurs for 5-fold validation, when the degree of temperature is 5. Although LOOCV is computationally more intense, here, 5-fold validation appears to show slightly better results. When applied to the whole dataset, the MSE values of all CV methods are almost similar.

3. Perform the analysis in problem no. 2, but this time, fit linear models with different combinations of X variables, without interaction. Discuss which model is most preferable based on the cross-validation results. Plot the results and on X-axis labels, provide the X-variable combinations used in the model, e.g. (temp, temp-press, temp-matfus, temp-matfus-mattr etc.)

To train linear models on different combinations of X, we have to create a formula list which enlists

all the different possible combinations among the five predictor variables. So, disregarding the null model, there will be a total of $2^5 - 1 = 31$ possible combinations

```
predictors <- c("Temperature", "Pressure", "Temp_x_Press", "MatFusion", "MatTransform")
model.formulas <- list()

# Generate all possible predictor combinations
for (i in 1:length(predictors)) {
  cmb <- combn(predictors, i, simplify = FALSE)
  for (combo in cmb) {
    formula_str <- paste("Quality ~", paste(combo, collapse = " + "))
    formula_obj <- as.formula(formula_str)
    model.formulas[[formula_str]] <- formula_obj
  }
}

n.models <- length(model.formulas)
```

Again, creating an empty list is created for storing errors and the models are trained.

```
cv.error.loocv <- rep(0, n.models)
cv.error.5fold <- rep(0, n.models)
cv.error.10fold <- rep(0, n.models)

model.names <- names(model.formulas)

for (i in 1:n.models) {
  start_time <- Sys.time()
  cat("Starting model", i, ":", model.names[i], "\n")

  glm.fit <- glm(model.formulas[[i]], data = data_sample)

  # LOOCV
  cv.loocv <- cv.glm(data_sample, glm.fit)
  cv.error.loocv[i] <- cv.loocv$delta[1]

  # 5-fold CV
  cv.5 <- cv.glm(data_sample, glm.fit, K = 5)
  cv.error.5fold[i] <- cv.5$delta[1]

  # 10-fold CV
  cv.10 <- cv.glm(data_sample, glm.fit, K = 10)
  cv.error.10fold[i] <- cv.10$delta[1]

  end_time <- Sys.time()
  duration <- round(as.numeric(difftime(end_time, start_time, units = "secs")), 2)
  cat("Finished model", i, "in", duration, "seconds. LOOCV MSE =", cv.error.loocv[i],
      "5-fold MSE =", cv.error.5fold[i], "10-fold MSE =", cv.error.10fold[i], "\n\n")
}
```

The first few and last few rows of training output are shown below.

```
Starting model 1 : Quality ~ Temperature
Finished model 1 in 16.39 seconds. LOOCV MSE = 133.0788 5-fold MSE = 133.0121 10-fold MSE = 133.0904

Starting model 2 : Quality ~ Pressure
Finished model 2 in 16.31 seconds. LOOCV MSE = 168.8959 5-fold MSE = 168.9318 10-fold MSE = 168.8616

Starting model 3 : Quality ~ Temp_x_Press
Finished model 3 in 17 seconds. LOOCV MSE = 157.702 5-fold MSE = 157.6551 10-fold MSE = 157.7967
```

.....

.....

```

Starting model 28 : Quality ~ Temperature + Pressure + MatFusion + MatTransform
Finished model 28 in 21.09 seconds. LOOCV MSE = 83.76925 5-fold MSE = 84.0066 10-fold MSE = 83.68985

Starting model 29 : Quality ~ Temperature + Temp_X_Press + MatFusion + MatTransform
Finished model 29 in 21.13 seconds. LOOCV MSE = 84.30565 5-fold MSE = 84.23364 10-fold MSE = 84.45012

Starting model 30 : Quality ~ Pressure + Temp_X_Press + MatFusion + MatTransform
Finished model 30 in 21.21 seconds. LOOCV MSE = 85.96859 5-fold MSE = 86.25406 10-fold MSE = 86.16316

Starting model 31 : Quality ~ Temperature + Pressure + Temp_X_Press + MatFusion + MatTransform
Finished model 31 in 23.48 seconds. LOOCV MSE = 83.83607 5-fold MSE = 84.00075 10-fold MSE = 83.79882

```

The error lists are converted to a dataframe and plotted.

```

combination.cv.results <- data.frame(
  Model = model.names,
  LOOCV_MSE = cv.error.loocv,
  CV5_MSE = cv.error.5fold,
  CV10_MSE = cv.error.10fold
)

[1] "Cross-Validation Errors for Different Predictor Combinations:"
> print(combination.cv.results)

```

	Model	LOOCV_MSE	CV5_MSE	CV10_MSE
1	Quality ~ Temperature	133.07880	133.01212	133.09036
2	Quality ~ Pressure	168.89589	168.93182	168.86156
3	Quality ~ Temp_X_Press	157.70200	157.65512	157.79669
4	Quality ~ MatFusion	124.77012	124.77034	124.65829
5	Quality ~ MatTransform	113.07841	113.24593	112.91921
6	Quality ~ Temperature + Pressure	133.14410	133.22270	133.21400
7	Quality ~ Temperature + Temp_X_Press	133.19586	133.62983	133.25670
8	Quality ~ Temperature + MatFusion	119.93438	119.98925	119.92527
9	Quality ~ Temperature + MatTransform	84.59484	84.73901	84.72504
10	Quality ~ Pressure + Temp_X_Press	138.69882	138.60187	138.70359
11	Quality ~ Pressure + MatFusion	123.41995	123.26544	123.45438
12	Quality ~ Pressure + MatTransform	113.13425	113.22824	113.07118
13	Quality ~ Temp_X_Press + MatFusion	121.81381	121.60052	121.76785
14	Quality ~ Temp_X_Press + MatTransform	112.26011	112.17882	112.12062
15	Quality ~ MatFusion + MatTransform	103.82232	103.60521	103.78660
16	Quality ~ Temperature + Pressure + Temp_X_Press	133.25959	133.94078	133.08580
17	Quality ~ Temperature + Pressure + MatFusion	98.12928	98.18716	98.37655
18	Quality ~ Temperature + Pressure + MatTransform	84.63338	84.71118	84.72145
19	Quality ~ Temperature + Temp_X_Press + MatFusion	101.81007	101.49119	101.76008
20	Quality ~ Temperature + Temp_X_Press + MatTransform	84.66181	84.48618	84.94776
21	Quality ~ Temperature + MatFusion + MatTransform	84.50945	84.45073	84.47078
22	Quality ~ Pressure + Temp_X_Press + MatFusion	120.03935	120.23471	120.20577
23	Quality ~ Pressure + Temp_X_Press + MatTransform	104.47502	104.51322	104.67427
24	Quality ~ Pressure + MatFusion + MatTransform	86.29017	86.48656	86.28865
25	Quality ~ Temp_X_Press + MatFusion + MatTransform	94.06184	94.69683	94.22870
26	Quality ~ Temperature + Pressure + Temp_X_Press + MatFusion	98.23669	98.21705	98.42299
27	Quality ~ Temperature + Pressure + Temp_X_Press + MatTransform	84.71167	84.44575	84.72333
28	Quality ~ Temperature + Pressure + MatFusion + MatTransform	83.76925	84.00660	83.68985
29	Quality ~ Temperature + Temp_X_Press + MatFusion + MatTransform	84.30565	84.23364	84.45012
30	Quality ~ Pressure + Temp_X_Press + MatFusion + MatTransform	85.96859	86.25406	86.16316
31	Quality ~ Temperature + Pressure + Temp_X_Press + MatFusion + MatTransform	83.83607	84.00075	83.79882

From different combinations, we check which combination gives the least MSE for any CV method applied.

```


> print(paste("Minimum MSE:", min_mse_value))
[1] "Minimum MSE: 83.6898466534485"
> print("Row(s) with the minimum MSE:")
[1] "Row(s) with the minimum MSE:"
> print(min_mse_row)

```

	Model	LOOCV_MSE	CV5_MSE	CV10_MSE
28	Quality ~ Temperature + Pressure + MatFusion + MatTransform	83.76925	84.0066	83.68985

The minimum MSE observed is **83.6898**. The predictor variables involved are **Temperature**, **Pressure**, **Material Fusion Metric** and **Material Transform Metric**. The minimum MSE was observed with the **10-fold validation method**.

```
cv.melted <- melt(combination.cv.results, id.vars = "Model", variable.name = "CV_Type", value.name = "MSE")
ggplot(cv.melted, aes(x = Model, y = MSE, fill = CV_Type)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Cross-Validation MSE for Different Predictor Combinations",
       x = "Predictor Combination",
       y = "MSE") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



4. Generate 50 random numbers from Normal Distribution $N(\mu = 50, \sigma^2 = 2)$. Now create 100 bootstrap samples with 20 data points each, with replacement. Estimate the mean and variance of the population from the bootstrap samples.

A new seed is set to generate 50 random numbers from the distribution $N(\mu = 50, \sigma^2 = 2)$.

```
> set.seed(789)
> pop_sample <- rnorm(50, mean = 50, sd = sqrt(2))
> head(pop_sample)
[1] 50.74118 46.80279 49.97217 50.25900 49.48897 49.31484
```

The bootstrap sample size is set as 20 and 100 such samples are generated after setting a new seed.

```
n_boot <- 100
boot_sample_size <- 20

boot.means <- rep(0, n_boot)
boot.vars <- rep(0, n_boot)

set.seed(456)

for (i in 1:n_boot) {
  boot.sample <- sample(pop_sample, size = boot_sample_size, replace = TRUE)
  boot.means[i] <- mean(boot.sample)
  boot.vars[i] <- var(boot.sample)
}

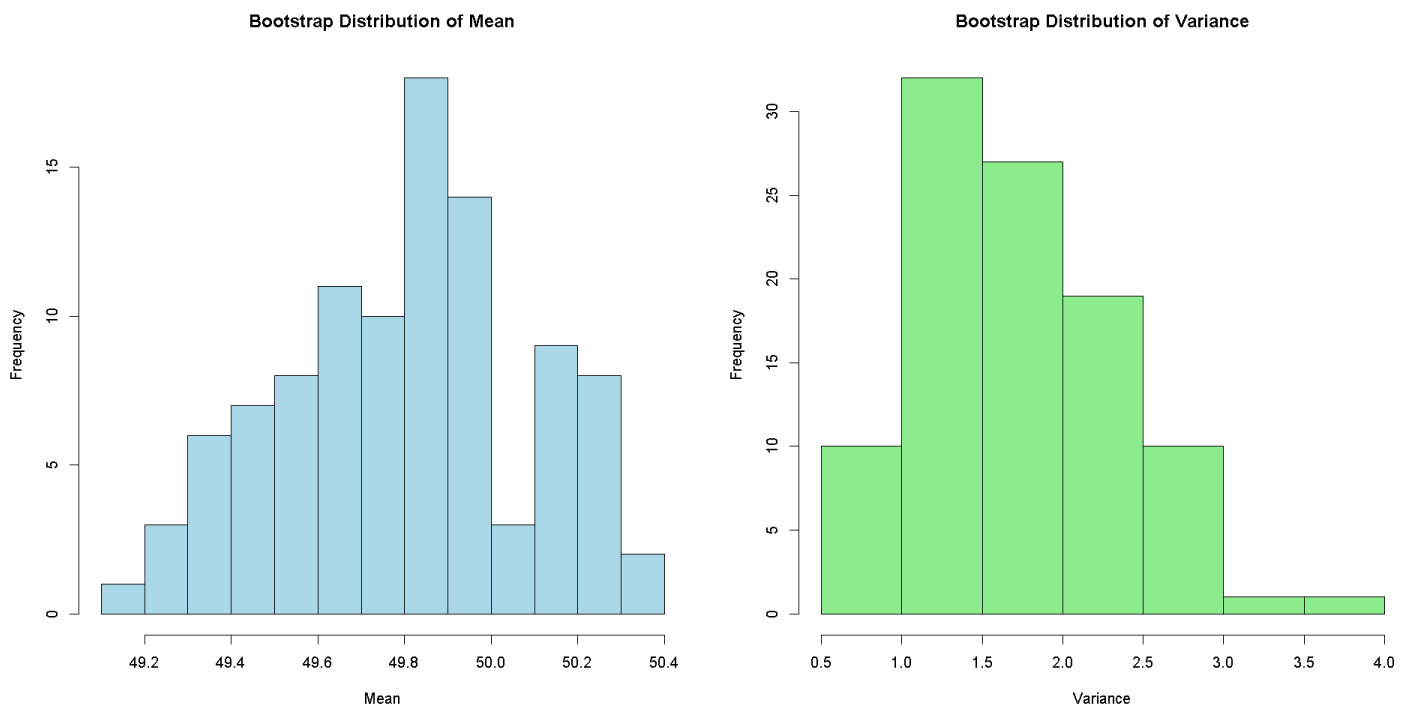
boot.mean.estimate <- mean(boot.means)
boot.var.estimate <- mean(boot.vars)
```

The estimated mean and variance from different samples are as follows.

```
> cat("Bootstrap Estimation Results (", n_boot, "samples of size", boot_sample_size, "):\n")
Bootstrap Estimation Results ( 100 samples of size 20 ):
> cat("Estimated Mean:", boot.mean.estimate, "\n")
Estimated Mean: 49.79905
> cat("Estimated Variance:", boot.var.estimate, "\n")
Estimated Variance: 1.699106
```

The frequency chart is plotted from the means and variance estimated from different samples.

```
par(mfrow = c(1,2))
hist(boot.means, col = "lightblue", main = "Bootstrap Distribution of Mean",
     xlab = "Mean", breaks = 10)
hist(boot.vars, col = "lightgreen", main = "Bootstrap Distribution of Variance",
     xlab = "Variance", breaks = 10)
par(mfrow = c(1,1))
```



Conclusions and Discussion:

- 5-fold validation yields lower MSE than LOOCV and 10-fold validation, even though LOOCV requires increased computational resources. Higher order polynomials tend to have lower MSE values across all CV methods, but have a risk of overfitting and loss of generalization.
- Among the 31 possible combinations of predictor variables, the model based on *Pressure*, *Temperature*, *Material Fusion Metric* and *Material Transform Metric*, gave the lowest MSE, using 10-fold technique. This shows that interaction effects and selection of suitable predictors affect the model performance. This might be an exception when 10-fold validation gives better results than LOOCV.

- Bootstrap resampling provided an estimate of the population mean and variance from 100 resampled datasets. The predicted variance was close to the actual variance ($\sigma^2 = 2$), showing the effectiveness in predicting population parameters.